VECTOR-SHIELD: SEMESTER 7 PROGRESS REPORT & SEMESTER 8 STRATEGIC ACTION PLAN

A Comprehensive Technical Review, Gap Analysis, Expanded Publication Strategy, and Enterprise Productization Roadmap

Institution:	Department of Artificial Intelligence, School of AI & Data Science, Amrita Vishwa Vidyapeetham, Coimbatore
AIDS Team 14:	Vignesh V S (cb.ai.u4aid23057) - Lead | Karthick Ruban T (cb.ai.u4aid23017) | Sanjay Senthil (cb.ai.u4aid23042) | Hari Prakash V (cb.ai.u4aid23047)
Supervisor & Guide:	**Dr. Lekshmi C. R.**, Assistant Professor, Department of Artificial Intelligence
Live Platform:	https://vectorshield.streamlit.app/ | GitHub: https://github.com/VigneshVS2005/Vectorshield
Academic Status:	Semester 7 Progress Review (Work Done So Far) | Target Technical Completion: Mid-December 2026

1. Foundational Concept: What Our Project is All About

The Problem: Enterprise AI agents are rapidly adopting Autonomous Retrieval-Augmented Generation (RAG) to execute critical operations—such as bank wire transfers, database modifications, and cloud API management. However, these systems are fatally vulnerable to **Indirect Prompt Injections (IPI)** and **Vector Database Poisoning**. Attackers conceal imperative natural language commands within external documents (PDF vendor invoices, regulatory updates). When retrieved by semantic search, these instructions hijack the LLM control flow, causing unauthorized actions.

Why Existing Tools Fail: State-of-the-art defenses (Meta's Llama Guard 3, NeMo Guardrails) inspect text *after* generation using secondary LLMs. This incurs **800–2,500 ms of latency**, massive continuous token costs, and aborts client sessions with destructive hard-crash errors.

The Vector-Shield Solution: Vector-Shield introduces an in-flight, dual-tier neuro-symbolic firewall directly in high-dimensional vector space ($\blacksquare^{384} \rightarrow \blacksquare^{16}$) *before* prompt synthesis. In Stage 1, non-linear scaled LOF catches obvious outliers in **1.82 ms with zero token cost**. In Stage 2, borderline ambiguous payloads are classified by a fine-tuned **149M ModernBERT** model in 12.2 ms. An autonomous **LangGraph self-healing loop** dynamically blacklists poison IDs and re-retrieves clean context, while automated **SIEM webhooks** notify enterprise Security Operations Centers (SOC).

2. Semester 7 Accomplishments So Far (Module-by-Module Audit)

Module	Subsystem / Artifact	Key Technical Engineering Highlights	Empirical Status
M1	Mathematical Topographic Metric	LOF scaled by centroid drift (d_cos/τ)^1.2 and 16-D adaptive PCA projection fitted on clean vectors.	**100% Outlier Recall**
M2	Ground-Truth Evaluation Benchmark	Verified corporate guidelines (SWIFT wire, SOX, GDPR) + real adversarial injections (Garak/HuggingFace).	**Verified Corpus**
M3	Stage 1 Symbolic Interceptor	Implemented in src/lof_interceptor.py; 384-D MiniLM encoding, sub-2ms PCA density filter.	**1.82 ms Latency**
M4	Stage 2 ModernBERT Classifier	Fine-tuned answerdotai/ModernBERT-base (149M params, 8,192 window, FlashAttention-2) on NVIDIA A100 GPU (FP16 AMP).	**100.00% Holdout F1**
M5	LangGraph Self-Healing Loop	Cyclic DAG (retrieve → intercept → heal → respond); dynamic blacklist rewrites; zero-crash execution.	**Zero Crash Execution**
M6	Enterprise SIEM Telemetry	Automated JSON alert dispatching to Splunk/Sentinel/HTTPBin; verified HTTP 200 OK telemetry.	**HTTP 200 OK Delivery**
M7	Production SaaS Dashboard	Deployed on Streamlit Cloud with 6 operational tabs, 3D PCA vector space, and multi-chunk PDF ingestion.	**Live Public SaaS**
M8	B.Tech Thesis LaTeX Draft	7 full chapters, IEEE/Amrita thesis formatting, empirical manifold diagrams, and test suites.	**7 Full Chapters**

3. Semester 7 Gap Analysis & Objective Self-Critique

To elevate Vector-Shield to premier conference publication tier and commercial enterprise readiness, Team 14 conducted an objective audit of current scope boundaries:

Gap ID	Current Scope Boundary	Technical Context	Next-Phase Strategic Solution
G1	**Dataset Breadth & Scale**	Focused initially on core banking/regulatory policies for architectural validation.	Scale to **10,000+ to 25,000+ samples** across 5 enterprise domains (Banking, Healthcare, Legal, DevOps, Code).
G2	**Modality Scope (Text Only)**	Focused on dense 384-D text representations.	Integrate **CLIP / SigLIP** multi-modal embeddings for visual prompt injections in scanned PDFs/charts.
G3	**In-Process Python Runtime**	Rapid prototyping in Python/Streamlit rather than compiled microservices.	Build an out-of-process **C++/Rust Envoy Sidecar** (<500µs latency) for Kubernetes pods.
G4	**Static Enterprise Centroid (c)**	Offline fitting on verified corporate documents.	Implement **Streaming Online Centroid Updates** (EMA) to adapt dynamically as documents evolve.
G5	**White-Box Gradient Attacks**	Tested empirical red-team sets (Garak) rather than direct gradient optimizers.	Stress-test against **Greedy Coordinate Gradient (GCG)** and HotFlip token optimization.
G6	**Single-Agent RAG Scope**	Evaluated on single-agent retrieve-and-generate workflows.	Expand LangGraph state machine to **Multi-Agent Swarms (AutoGen/CrewAI)** consensus guardrails.

4. Immediate Strategic Action Plan (Target Completion: Mid-December 2026)

Our agile action plan targets completing core technical enhancements by **mid-December 2026**, transitioning immediately to peer-reviewed paper submissions and commercial productization:

Action Item	Key Technical Engineering Objective	Industrial & Academic Impact
Item 1: Massive Dataset Scaling	Scale benchmark corpus to **10,000+ to 25,000+ authentic samples** across 5 domains: Banking (ISO 20022), Healthcare (HIPAA), Legal (SEC EDGAR), Cloud DevOps (AWS IAM), Code RAG.	Guarantees broad cross-domain enterprise compatibility and high statistical power for peer review.
Item 2: Multi-Modal Topographic Shield	Integrate **CLIP ViT-L/14 & SigLIP** multi-modal embeddings to intercept visual prompt injections in scanned PDF invoices, receipts, and infographics.	First-of-its-kind visual injection firewall for enterprise multimodal RAG systems.
Item 3: C++/Rust Envoy Sidecar	Compile PCA-LOF into an out-of-process **C++/Rust Envoy proxy filter** (<500µs latency, Kubernetes Helm chart).	Enables zero-code drop-in defense for any enterprise microservice running in Kubernetes.
Item 4: White-Box Adversarial Hardening	Subject system to **Greedy Coordinate Gradient (GCG)** and HotFlip attacks; train adversarial defense in ModernBERT.	Proves mathematical resilience against sophisticated mathematical token optimizers.
Item 5: Turnkey Vector DB Connectors	Native plugins for **Pinecone, Milvus, Qdrant, OpenSearch** + streaming online centroid updates (EMA).	Seamless enterprise adoption with zero manual re-indexing required as knowledge grows.
Item 6: Multi-Agent Swarm Defense	Consensus-driven guardrail for **AutoGen / CrewAI** swarms to prevent cascading agent-to-agent infection.	Protects next-generation collaborative AI agent architectures from systemic collapse.

5. Expanded Research Publication Strategy & Target Venues

Proposed Research Paper Title: *'Topographic Anomaly Detection and Stateful Semantic Guardrails in High-Dimensional Embedding Spaces for Securing Autonomous Agentic RAG'*
Expanded Target Venues:
• **Tier-1 Cybersecurity:** USENIX Security, ACM CCS, IEEE S&P; (Oakland), NDSS.
• **Tier-1 AI / NLP:** ACL, EMNLP, NAACL, NeurIPS / ICLR (Safety & Alignment Tracks).
• **Premier Flagship Journals:** IEEE Transactions on Information Forensics and Security (TIFS), IEEE Transactions on Dependable and Secure Computing (TDSC), Computers & Security (COSE).
Benchmark Baselines for Paper Acceptance: Llama Guard 3, NeMo Guardrails, SmoothLLM, PromptGuard-86M, Securita.
Intellectual Property & Patents: Indian Provisional Patent Application covering: (1) Non-linear centroid-drift scaled LOF in PCA-reduced embedding manifolds; (2) Cyclic state-machine self-healing for agentic workflows.

6. Maximum Product Vision & Commercialization Blueprint

Maximum Product Vision (What Vector-Shield Can Achieve at Its Peak):
1. **Universal Enterprise Security Mesh:** Lightweight C++/Rust Envoy sidecar intercepting every vector DB call in <500µs with zero application code changes.
2. **Multi-Modal Ingestion Firewall:** Comprehensive screening of both text chunks and scanned PDF/image invoices before persistent database indexing.
3. **Self-Healing Agentic Execution:** Dynamic LangGraph state recovery guaranteeing agents never crash or refuse legitimate users during attacks.
4. **Multi-Agent Consensus Guardrail:** Eliminating cascading infections across collaborative agent swarms (AutoGen, CrewAI).
5. **Turnkey Vector DB Connectors:** One-click integration with Pinecone, Milvus, Qdrant, and AWS OpenSearch.
6. **Automated SOC Compliance:** Structured JSON telemetry for SOC 2 Type II, ISO 27001, and HIPAA compliance audits.
Immediate Steps Leading to Mid-December: Finalize 10k–25k scaling, complete multi-modal screening, release arXiv preprint, and package Kubernetes Helm chart for enterprise pilots.

7. Formal Supervisor Endorsement Request

AIDS Team 14 respectfully submits this progress report to our project supervisor, **Dr. Lekshmi C. R.** Semester 7 deliverables have met all departmental requirements with high empirical rigor. We request formal approval to proceed with this strategic action plan toward mid-December completion.

Vignesh V S
(Lead Architect, AIDS Team 14)

Dr. Lekshmi C. R.
(Project Supervisor & Advisor)

Department of AI
Amrita Vishwa Vidyapeetham